REPORT FOR 1DV516

ALGORITHMS AND ADVANCED DATA STRUCTURES

INGMAR IMMANUEL FALK

SEPTEMBER 17, 2023

Contents

1 Abstract

This report will go over how I solved the problems in the assignment.

2 Introduction

In this report I will elaborate my solutions to the required problems (4-7) as well as the problems for the grade VG (8). I will start with a general description of the setup for the experiments, how I tested my algorithms for both correctness and speed and how I then used those results to create visualizations of the data.

3 Experiment Setup

Before I will go into the details of the testing, benchmarking and visualizations. I will elaborate the general structure of the project and the code.

3.1 Structure

I have structured the project so that every major responsibility is contained within its own folder.

`bencharks` contains all benchmarking code

`bench_results` contains all benchmarking results as *.csv* files

`tests` contains all the testing code

`analysis` contains all the python code for data analysis

`fits` contains all the plots with fitted graphs generated using *scipy*

`calcs` contains all the results for the fits generated using *scipy*

`plots` contains all the generated visualizations as *.png* files

`bin` contains the compiled java code

`src` contains all the source code for the algorithms as well as the helper utility for the testing and benchmarking

Running any functionality of the project should be done via *make*. The *Makefile* is located in the root of the project and contains all the necessary commands to build the java code, run tests and benchmarks and generate the visualizations.

3.2 Testing

To test the correctness of the algorithms I have created a couple of helper classes that allow me to string together a series of tests and then run them. The results are printed to *stdout* and contain information about which tests passed and which failed. The Makefile contains commands to run all tests together or just a single file. Its not very advanced but it is enough for the purpose of this assignment.

My tests were simple, I took small inputs for which I could figure out the supposed output myself and used it to test the algorithms. For some algorithms I added more tests, just to be sure, because while the algorithms are fairly trivial to implement, the improved cache version required slightly more thinking. But I will talk more about that later in Problem 5 and Problem 6.

3.3 Benchmarking

For the benchmarking I did something similar and created a set of utility classes that use Java's *System.nanoTime()* to measure the time it takes to run a given algorithm. To make the *bench* function I used for all my tests generic, I this public interface:

```
public <T> void bench(
    String name,
    String filename,
    Function<Integer, T> setup,
    Consumer<T> fn,
    List<Integer> sizes,
    int reps
) { ... }
```

The parameters *name* and *filename* are quite trivial and simply used to display the name of the benchmark to *stdout* and to save the results to a *.csv* file. The rest requires some more explanation:

- *setup* takes the size to test an arbitrary algorithm at as input and returns the algorithms input as type *T*. This is so that the setup can be adjusted according to the size. It is also crucial to our ability to make the benchmarks multithreadable. This setup function allows us

to isolate each thread and encapsulate everything it needs so we dont have any shared state.

- *fn* is a function that takes the input generated by the *setup* of type *T* as input and returns nothing. This function is the algorithm that is being benchmarked.

- *sizes* is a list of integers that represent the input sizes that should be benchmarked. The benchmarking function will run the algorithm for each input size and measure the time it takes to run the algorithm for each input size.

- *reps* is the number of times the algorithm should be run for each input size. The benchmarking function will run the algorithm *reps* times for each input size and then calculate the average time it took to run the algorithm for each input size.

While we get the time as *nano seconds*, we store and display it in *milli seconds* to make it easier to read.

3.4 Visualizations

For the visualizations I used *python* and *matplotlib*, *numpy* and *scipy*. There is not much to say here. I simply load the data using numpy and then pass use it with matplotlib to generate simple graphs. For the fitting I used scipy's *curve_fit* function together with a couple of simple functions I defined such as linear, quadratic, cubic etc..

4 Problems - G

4.1 Problem 4 - UnionFind Benchmarking & Growth

For the comparison between the basic Quick Find version of Union Find and some improved version, I chose the Weighted Quick Union *(WQU)*. And although I also implemented all other options, I will only use the results for the *WQU* implementation for the comparison.

4.1.1 Quick Union Expectations

Before looking at the benchmarks and calculations we need to think about what we expect to be able to check if we messed up our implementation or if it was correct.

The dominant factor when to consider when arguing about the runtime performance of the basic quick find implementation is the *union* function and here in particular the *for loop* over all N elements. This will obviously result in a time complexity of *O(N)* as other operations such as retreiving the id/root are *O(1)* operations and since we ignore lower order terms and coefficients that leaves us with *O(N)*.

4.1.2 Weighted Quick Union Expectations

For the *WQU* however we expect something quite a bit better than *O(N)*. The reason for this is the check we perform in every union. We maintain a list of weights for all nodes, and with every union, the weight of the two nodes to be connected will be compared and the node that has a lower weight will be appended to the heavier one. I am using weight instead of size as this name more accurately describes what we are keeping track of here. Since we are only keeping track of the total number of nodes in one tree, not the depth/size of it. For a list of length N we could have a single union with weight N but a max depth of 1.

Now the reason this check will give us a huge performance boost is because now, the depth of the tree can only be increased when two unions of the same size are connected together. This means that in order to get the max depth of the tree, we would have to only ever unionise nodes with the same weight, which is equivalent to doubling the weight of all nodes before moving on to the next round an doing the same thing again. Lets look at an example:

Lets start with a newly initialised list of size 8. We have 8 root nodes with a weight of 1 [(0, 1), (1, 1), (2, 1), (3, 1), (4, 1), (5, 1), (6, 1), (7, 1)]. In the first round, we unionise 4 unique pairs. This will result in 4 root nodes with a weight of 2 [(1, 1), (1, 2), (3, 1), (3, 2), (5, 1), (5, 2), (7, 1), (7, 2)]. By this point, the max depth is 1. The next iteration will result in 2 root nodes with a weight of 4 and a max depth of 2 and the last iteration last in one root node with a weight of 8 and a max depth of 3. At this point its weight is equal to N, since all of the nodes within the list have it as its root node.

If we want to apply this to any generic list of size N, we have to think about how many of those iterations can we possibly do. This is the same as asking, how many times can we *half* the number of N root nodes or how many times we can double 1. Both of these will lead to the same result, which is $log_2 N$. This holds not only for lengths of a power of 2, but for any N, since in those cases the depth is still dominated by the largest power of 2 that is smaller than N.

All of this should lead uf to the expectation that the *find* or *root* function of the Weighted Quick Union Find should have a time complexity of *O(log N)*. If we think about the time complexity of the *union* and *connected* function, where we have to call find/root twice, it might seem like the time complexity should be *O(2log N)* but since we ignore coefficients this will still result in *O(log N)* for both functions.

4.1.3 Quick Union Results

The benchmarking for the basic quick find implementation was done by running the *union* function on a list of N sizes ranging from 10.000 to 1.000.000 in steps of 10.000. Each function will be called 1.000 times. The resulting growth we can observe is as expected linear. To facilitate this lets take a couple of points to calculate the necessary values. For this

size	time	l2 t	ratio	l2 r	slope	b	c
100000	15	3.95	N/A	N/A	N/A	N/A	N/A
200000	27	4.78	1.77	0.83	0.000120	0.826	-9.76
300000	44	5.49	1.63	0.71	0.000170	1.208	-16.50
400000	59	5.89	1.32	0.40	0.000145	0.972	-12.20
500000	72	6.19	1.23	0.30	0.000136	0.923	-11.28
600000	87	6.45	1.20	0.26	0.000144	0.991	-12.58
700000	104	6.71	1.20	0.26	0.000175	1.188	-16.36
800000	114	6.85	1.10	0.13	0.000102	0.695	-6.78
900000	132	7.05	1.15	0.20	0.000171	1.179	-16.28
average		6.17	1.33	0.39	0.000146	0.998	-12.72

Table 1: All the results of calculations based on the *size* and *time*. *l2 t* = $log_2 time$, *l2 r* = $log_2 ratio$

The *size* and *time* column are quite obvious to understand. The $log_2 t$ column contains simply the log_2 of the original time values for comparing

the values we get from our calculated function and confirming its correctness. The same goes for the *ratio/log_2ratio* columns. To calculate the slope we take the size and time of two rows and put them in this formula: $\frac{time1-time2}{size1-size2}$. The log_2slope is basically the same but instead of taking the raw values we use the log_2 of them, like this $\frac{log_2time1-log_2time2}{log_2size1-log_2size2}$. Once we have the b, we can use it to calculate the constant c like this: $c = log_2time - b * log_2size$. The reason why there are no values in the first row is that since we rely on the previous value for all results and there are no previous values for the first row, we are unable to calculate anything.

But now that we calculated the b and c for some points, we can take the average and use that to create a mode that predicts/approximates the *log_2 of a time* for a given size. The function in question looks like this $log_2t = b * log_2n + c$, where t is the time and n the size. In our case, the average b is *0.998* and the average c is *-12.72*, which leaves us with the following function: $log_2t = 0.998 * log_2n - 12.72$. Lets compare the values we get from the model predictions to the original values.

size	original log_2t	prediction	accuracy
200000	4.78	4.85	0.98
300000	5.49	5.44	1.01
400000	5.89	5.85	1.01
500000	6.19	6.17	1.00
600000	6.45	6.43	1.00
700000	6.71	6.66	1.01
800000	6.85	6.85	1.00
900000	7.05	7.02	1.00

Table 2: Model Predictions vs Original log_2 times calculated with b and c

As we can see, the model predictions are quite accurate, with only minimal errors in their accuracy. But these are only the log_2 values, to get the actual time we have to create a power law that approximates the actual time rather than just the log_2 of it. This is done by simply taking the already calculated b and c values and plugging them into the general power law formula $a * x^b$, where $a = 2^c$. This results in the following function: $t = 2^{-12.72} * n^{0.998}$ or $t = 0.00015 * n^{0.998}$. Lets compare the values we get from the model predictions to the original values:

Once again, the predicted values are fairly close to the original times,

size	original t	prediction	accuracy
100000	15.49	14.47	1.07
200000	27.45	28.90	0.95
300000	44.81	43.31	1.03
400000	59.27	57.71	1.03
500000	72.83	72.11	1.01
600000	87.25	86.49	1.01
700000	104.79	100.88	1.04
800000	114.98	115.25	1.00
900000	132.11	129.63	1.02

Table 3: Power Law predictions vs Original times calculated with b and c

which indicates that our calculated values are indeed correct. To fully confirm this, I also used *scipys curve_fit* function to generate a power law in order to compare it to our calculated one. Important to mention here is the fact that the curve_fit was ran on all of the data, not just these 9 values. The resulting power law is the following

a = 0.00016, b = 0.994 vs our results: a = 0.00015, b = 0.998

We can confidently say that our calculations were accurate. That now allows us to argue about wether our expectations about the growth rate of the quick find were correct. The calculated exponents (b) are very close to 1 which indicates linear growth and confirms our assumption that the basic implementation of union find are linear. This first analysis was a bit more in depth than the following ones since all the rules and formulas had to be established.

4.1.4   Weighted Quick Union Results

Since now the process is established we can go straight into the calculations and their results. We apply the same steps as previously, take the sizes and times, calculate the slopes, b's and c's and use it to create the general function to predict the $log_2 t$ and the power law to approximate the actual size. This leaves us with the following tables.

Using the *b (2.117)* and *c (-13.63)* to create a power law gives us the following: a = 7.88e-5 = 0.0000788, b = 2.117

size	time	l2 t	ratio	l2 r	slope	b	c
100000	1.94	0.95	N/A	N/A	N/A	N/A	N/A
200000	2.03	1.02	1.05	0.07	9.1965e-07	0.067	-0.16
300000	1.95	0.97	0.96	-0.05	-7.3533e-07	-0.091	2.62
400000	2.29	1.20	1.17	0.23	3.3610e-06	0.551	-9.07
500000	1.73	0.79	0.76	-0.40	-5.5719e-06	-1.249	24.44
600000	1.96	0.97	1.13	0.18	2.2789e-06	0.677	-12.03
700000	2.03	1.02	1.04	0.05	7.2662e-07	0.236	-3.56
800000	1.90	0.93	0.94	-0.09	-1.2665e-06	-0.481	10.37
900000	1.98	0.99	1.04	0.05	7.2805e-07	0.318	-5.30
average		0.99	1.01	0.00	5.51e-08	0.004	0.91

[h]

size	original	pred	acc	size	original	pred	acc
100000	1.9364	1.9637	0.9861	100000	0.9534	0.9735	0.9793
200000	2.0283	1.9685	1.0304	200000	1.0203	0.9771	1.0442
300000	1.9548	1.9713	0.9916	300000	0.9670	0.9792	0.9876
400000	2.2909	1.9733	1.1609	400000	1.1959	0.9806	1.2195
500000	1.7337	1.9749	0.8779	500000	0.7939	0.9818	0.8086
600000	1.9616	1.9762	0.9926	600000	0.9720	0.9827	0.9891
700000	2.0343	1.9773	1.0288	700000	1.0245	0.9835	1.0417
800000	1.9076	1.9782	0.9643	800000	0.9318	0.9842	0.9467
900000	1.9804	1.9790	1.0007	900000	0.9858	0.9848	1.0010

Table 4: Power Law Predictions vs Original Time

Table 5: Model Predictions vs Original $log_2 Time$

As we can see in Table **??**, the results are drastically different from the basic quick find implementation. The slope for example is at some points negative and adds up to an average of near *0* (0.0000000551). The calculated b and c values seem to check out as the predictions calculated with it are only slightly off from the original ones. If we look at the *scipy* generated power law however, it looks quite a lot different.

a: 3.145, b: -0.034 vs a: 1.879, b: 0.004

or

$y = 3.145 * x^{-0.034}$ vs $1.879 * x^{0.004}$

Now those are quite different results. This is most likely due to the fact that for the *scipy* version I again used all of the data, rather than just this subset, but it is interesting to see, that such different values can be used to

get bascially the same output. Nevertheless, we can use our calculations to evaluate our predictions. Even tho both our calculations and the generated ones are quite different, they have in common the fact that they are both *sub-linear*. This means that as the input size grows, the slope/growth rate decreases until. If we look at the recorded times we can see, not only does the growth rate slow down, it is nearly constant over all 9 observed instances while fluctuating around 2. The average ratio indicates *no/minimal* growth. This indicates similar behaviour in terms of growth rate to the logartihmic ones and is consistent with our expectation of the worst case time complexity being *O(logN)*.

4.2   Problem 5 - Threesum Brute Force

The implementation of the brute-force threesum was straight forward. Simply nest three for loops and check if the sum of the three numbers is zero in the third loop. In order to avoid using the same number in the calculation, we check the indicies of the numbers we are currently checking and just continue the loop if any of the indicies of the loops match. To keep track of the three numbers I created a simple *record* that stores the three numbers. This will include duplicates, meaning for example that *(1, -2, -6, 1)* will result in *(1, -2, 1), (1, 1, -2), (-2, 1, 1), (-2, 1, 1), (1, -2, 1), (1, 1, -2)* for a total of *3!* results for a single valid triple. This includes results that are the same simply reordered versions of an already existing one.

4.3   Problem 6 - Threesum Improved

The siutation with the improved three sum was a bit weird for me. Originally I implemented the cached version as the optimisation. After a bit of thinking I thought of a way to implement a cached version that while being $O(N^2/2)$, so $O(N^2)$, runs nearly linear for all sizes if the range of the numbers contained is small (e.g. (-10, 10)). The issue with this was with the programming language that we have to use being java. When I ran my run my tests using I get a linear growth until 16.000, after which there is extreme fluctuations in the time. This is not due to the algorithm, as I implemented the same version in rust, to see if I had the same issue there, but low and behold, it runs linearly up to 50.000 extremely steady with no fluctuations. That just shows that the issue is most likely in java.

Figure 1: Java implementation up to sizes of 50.000

Figure 2: Rust implementation up to sizes of 50.000

Due to this reason, I decided that in this report I will go with the *two pointers* version, since java is delivering consistent data for this. But as a downside, it is of course way worse when it comes to growth rates so I will only be able to run the benchmarks with sizes up to 5.000 in steps of 100 with 20 repetitions, otherwise these would take way to long. To my discontent, the two pointer version is as bad as the brute force version when it comes to duplication of already found triples, something that was often halfed in the improved cached version, without duplication checking. I did not add any checking for duplicates in any implementation since that would just mean an increase in runtime and possibly scewing of the benchmarks.

4.3.1   Two Pointers

The two pointer implementation was the most verbose out of all of the algorithms. It is based on two pointers that are moving from opposite sites of the list and eventually meet somewhere in the list. It is required that the list is sorted for this to work, which adds more complexity and time to the algorithm. Sorting a list, if a proper algorithm is used, has an $O(NlogN)$ time complexity. The outer for loop and the inner while loop combined with $O(N)$ each will combined result in $O(N^2)$. The all the operations within the loops, such as if statements and array indexing can be ignored. This will leave us with $O(NlogN) + O(N^2)$, but since $O(N^2)$ dominates in this case, the total time complexity of this algorithm will be $O(N^2)$.

4.4 Problem 7 - Threesum Benchmarking & Growth

As mentioned in the preivous section, for this algorithm, the maximum sizes we are able to test it at are significantly lower than for the union find. So in this case, the tests were run for ranges from 100 to 5000 with steps of 100. As samples, I took 1, 2, 3 and 4 thousand. And to my pleasant surprise, even just taking such a small amount of values was enough for quite accurate values.

Table 6: All the results of calculations based on the *size* and *time*. $l2\ t = log_2 time$, $l2\ r = log_2 ratio$

size	time	l2t	ratio	l2r	slope	b	c
1000	186.63	7.54	N/A	N/A	N/A	N/A	N/A
2000	749.74	9.55	4.02	2.01	5.6311e-01	2.006	-12.45
3000	1864.80	10.86	2.49	1.31	1.1151e+00	2.247	-15.09
4000	3408.16	11.73	1.83	0.87	1.5434e+00	2.096	-13.35
average		10.72	2.78	1.40	1.07e+00	2.117	-13.63

Using the *b (2.117)* and *c (-13.63)* to create the power law results in the following. a = 7.89e-5 = 0.0000789, b = 2.117, power law: $7.889e - 05x^2.117$. We can then again use those numbers to approximate the time for any size and get fairly accurate results.

size	original	pred	acc
1000	186.63	176.46	1.06
2000	749.74	765.21	0.98
3000	1864.80	1805.03	1.03
4000	3408.16	3318.35	1.03

size	original	pred	acc
1000	7.54	7.46	1.01
2000	9.55	9.58	0.99
3000	10.86	10.82	1.00
4000	11.73	11.70	1.00

Table 7: power law predictions vs original time

Table 8: model predictions vs original $log_2 time$

To determine the growth rate of the two pointers algorithm we can look at the *b* value. In this case it is around 2.1. While the quick find around 1 indicates linear growth and the sub 1 *b* of the weighted quick union implementation indicated logarithmic/sublinear rates, an exponent b of two suggests a quadratic growth, meaning the bigger the input size, the steeper the slope. This however confirms our expectation of a quadratic growth rate and therefore should qualify as a threesum that has an upper bound of $O(N^2)$.

13

In order to ensure that the the two pointer implementation actually has an upper bound of $O(N^2)$ I compared it to both my original cached version, which consists of two for loops and keeps track of a set, therefore being $O(N^2)$ and the improved cached version, both of which have a time complexity of $O(N^2)$.


Figure 3: Comparison between implementations

As can be observed in the diagram, the two pointer implementation lies right in between the two cache implementations, leading us to the confirmation that it has both a $\Omega(N^2)$ as well as an $O(N^2)$.

5   Problems - VG

5.1   Problem 8 - Percolation

In this section I will quickly go over how I implemented problem 8 and then tested it as well as my approach for determining p^*.

5.1.1   Implementation

The main two components of my implementation are a n * n zeroed grid (int[]) and one of the union find implementations, in my case the weighted quick union find with path compression. As already alluded to in the task description, when instantiating the union find object, instead of initiating it

just with n^2, we add two more sites to it, so called virtual sites. One will serve as the top and the other as the bottom site. Meaning that all the grid values are being stored from 0 to n - 1 and the top value will be n^2 and the bottom $n^2 + 1$.

The two crucial methods will be the *open* and *percolates* method. While the *percolates* method is simply calling the union finds *connected* function with the top and bottom value to check if there is a valid path through the grid, the open method is a bit more complex. After validating that the row and column are indeed valid to be inserted and the site they describe is not already open, we open the site by setting it to 1 and increasing the internal tracker of how many sites are already open. In order to connect new opened sites now, we check all fields horizontally and vertically adjacent to it. If the site is also open, we call the *union* method of the union find on both points.

5.1.2 Testing

In order to make sure that the implementation was working, I created a simple *4x4* grid and opened specific fields manually in a way that I knew they had to make the grid percolating. This confirmed that what I had done worked correctly so I could move on to calculating the threshold.

5.1.3 Calculating threshold

For the calculation of the p threshold I, as suggested, made use of the *monte carlo* approach. This quite frankly means, taking randomly generated input to approximate a value, that could have been deterministically evaluated, but might take a lot of effort and time. In our case, that means initialising an *Percolation* instance for a size n, and randomly opening a site until the grid percolates. The number of trials in this case is simply n^2. This is done multiple times, in this case 500 times, after which the average is taken and those results used to calculate the standard deviation. This was done for 250 n's ranging from 2 to 500. This resulted in a final result of *p* = 0.59191* with a standard deviation of *std_dev = 0.01261*.